

B-cell epitope prediction from antigen-only dynamics: project log

Brandon Neff

John Altin

TGen

July 1, 2026

Living record. Read top to bottom to get up to speed.

Running record of the project: predict B-cell epitope residues from *antigen-only* information, using alanine-scan $\Delta\Delta G_{\text{bind}}$ as residue-level binding-importance labels and asking whether antigen dynamics / solvation / energetics carry predictive signal *beyond static surface exposure*. Read in order; each section is one iteration — what we tried, what we found, what it changed. Not a formal manuscript. Every figure is regenerated by a named script (recorded in **figures/FIGURES.md**); every number traces to a file or command; unconfirmed values are flagged [verify].

1 Motivation and the gap

We want to predict which residues of an antigen are B-cell epitope residues, from the antigen structure alone, in cases where the antibody is unknown. Doing this rigorously requires residue-level labels for “binding importance.” We take these from **alanine scanning**: for residue i , $\Delta\Delta G_{\text{bind}}^{(i)} = \Delta G_{\text{bind}}^{\text{mut},i} - \Delta G_{\text{bind}}^{\text{WT}}$, with large positive values flagging residues whose wild-type identity stabilizes the interface. We treat this as a residue-level epitope score, $S_i^{\text{epitope}} \approx \Delta\Delta G_{\text{bind}}^{(i)}$. (Full derivation and the alchemical / StaB-ddG label tiers: `../tcr/docs/b_cell_eptiope_prediction_may23_update_final.pdf`)

The gap. Most B-cell epitope predictors — and the benchmarks that train them — label residues by *contact distance* to a bound antibody. As a coarse first pass this is defensible: proximity to the paratope is a reasonable way to *nominate* candidate interface residues. The limitation is treating every contact as a positive label of equal weight, which conflates “close” with “important”: many interface residues contribute little binding free energy, while a few hotspots carry most of the affinity. The bias is baked into the data itself. Of the 4,164 conformational (discontinuous) B-cell epitopes in IEDB [1] that carry a deposited structure, essentially all (4,161; > 99%) are defined by crystallographic or cryo-EM *contact* rather than by any binding measurement;¹ models trained on these labels therefore inherit proximity as their working definition of “epitope.” Alanine-scan $\Delta\Delta G$ supplies the missing energetic refinement — *which* of the contacting residues actually pay for binding. Even a binary critical/non-critical call from a saturating alanine scan is energetically grounded, since “critical” means mutation measurably weakens binding; proximity and energetics are thus complementary, the former a candidate filter and the latter the importance label.

Separately, since Westhof et al. [2] it has been known that antigenicity correlates with *segmental mobility* (flexibility / B-factor) — an antigen-intrinsic, antibody-agnostic property — and

¹IEDB next-generation query API (query-api.iedb.org), `bcell_search`, structure type *Discontinuous peptide* with a non-null `pdb_id`, queried 2026-06. IEDB attaches a structure only when the epitope was itself defined from one, so structure-linked records are overwhelmingly contact-derived; the mutation/binding-derived epitopes carry no `pdb_id`.

energy-decomposition analyses of plain MD trajectories can already localize antibody epitopes on an antigen without knowledge of the antibody [3]. The opportunity is to test, with modern MD, whether antigen-only *dynamical* and *solvation* features (RMSF, collective-mode participation, hydration thermodynamics) predict energetic binding importance — and whether they beat static exposure. This is where the author’s methods background (FRESEAN collective-mode analysis; 3D-2PT-style hydration thermodynamics; non-equilibrium energy flow) is the differentiator (`../../_literature/my_manuscripts/`).

2 Design decisions to hold onto

Five issues decide whether the project produces a result or an artifact. Recorded here so they stay front-of-mind.

1. **The label is antibody-specific; the features are antibody-agnostic.** $\Delta\Delta G_{\text{bind}}^{(i)}$ is a residue’s contribution to *one* paratope, but $\mathbf{x}_i$ comes from the apo antigen and cannot know the antibody. So $f(\mathbf{x}_i) \rightarrow \Delta\Delta G^{(i)}$ can only recover the *antibody-marginal* (antigen-intrinsic) hotspot propensity. This is not a weakness to hide — it is exactly why dynamics/flexibility features should help (cf. Westhof). Where an antigen has several antibodies, average $\Delta\Delta G^{(i)}$ over them to *define* the intrinsic target.
2. **“Beyond SASA” is the primary hypothesis, not a footnote.** Surface exposure is the master confounder: exposed residues are simultaneously more likely to be hotspots, more flexible, and more hydrated. The publishable claim is not “flexibility predicts epitopes” (Westhof has that) but “dynamics/solvation features add signal *after* controlling for static exposure.” Design: baseline = SASA + depth + residue type; then test the partial R^2 added by RMSF, mode participation, and 3D-2PT hydration.
3. **Effective $N \approx \#\text{antigens/folds}$, not $\#\text{residues}$.** Residues within an antigen are correlated and the intrinsic signal is per-fold. Split **leave-one-complex-out** (ideally leave-one-fold-out); never random residue splits. At this scale the work is hypothesis-generating — powered for a strong dynamics-vs-exposure effect, not a deployable predictor. State this plainly.
4. **Features come from the apo antigen** — which is the point and sets a ceiling. Cryptic / induced-fit epitopes (conformational selection vs. induced fit) cannot be predicted in principle from apo features. Conformational-selection epitopes *should* show in apo flexibility/mode features; pure induced-fit ones will not. That contrast is itself a result.
5. **The MD protocol is not one stack.** A flexibility/RMSF run differs from a FRESEAN/3D-2PT run (the latter needs high-frequency velocity output, constraint removal, and a water model validated for dynamics). Likely two protocols per antigen. The house TIP3P/2 fs/LINCS default is scrutinized accordingly (see `../docs/METHODS.md`).

3 The label benchmark

Curated experimental alanine-scan $\Delta\Delta G$ for antigen-side residues lives in `benchmark/antigen_alanine_scanning` (6 sheets: README, Simple_Primary, Full_Schema [51 cols], hGH_Table2_4, Paper_Inventory, Counts). Audited 2026-06-18.

Coverage. 203 primary rows, 13 datasets, 7 antigens. Source: AB-Bind [4] (195 rows; github.com/sarahsirin/) and the Jin hGH alanine scan [5] (8 rows; benchmark cites the 1994 thesis, Table 3.2). The dataset was subsequently expanded with three antigen-side scans from SKEMPI 2.0 [6] and a binary label tier from IEDB (§5 and the Tier-2 section). Sign convention: $+\Delta\Delta G \Rightarrow$ weaker binding $\Rightarrow$ residue important (38 negative = binding-improving alanines; 17 zeros).

Integrity. The $\Delta\Delta G \leftrightarrow K_d$ -ratio arithmetic is internally consistent for all 203 rows ($< 2\%$ vs $\exp(\Delta\Delta G/RT)$). No censoring in the primary set. The 4 non-alanine rows (hGH, EC50-derived) are correctly flagged `exclude_from_binding_label`. Provenance (DOI, source line, URL) is complete except the two homology-modeled BoNT/A2 sets and the Jin rows. `transcription_confidence` is uniformly **high** — treat as a soft flag, possibly a default fill. [verify] spot-check 1–2 known hotspots against the source papers.

Clean MD-correlation core. Dropping homology models (HM_) and EC50/excluded rows leaves **163 rows / 5 antigens on real PDBs** (VEGF, lysozyme, HPr, BoNT/A1, MT-SP1).

4 First system: hen egg-white lysozyme (HEL)

HEL is the cleanest entry point: small (129 residues), rigid (apo $\approx$ bound), very stable, textbook AMBER behavior, and — crucially — it carries **three** independent antibody alanine scans in the benchmark, on a consistent 1–129 numbering:

Antibody	PDB	Ag chain	# Ala	$\Delta\Delta G$ range (kcal/mol)
HyHEL-63	1DQJ	C	11	0.60–6.10
D1.3	1VFB	C	12	0.20–2.90
HyHEL-10	3HFM	Y	4	1.03–6.83

Why this is better than one antibody. HyHEL-10 and HyHEL-63 share an epitope and *agree* on hotspots (K96: 6.83 vs 6.10; Y20: 4.87 vs 3.30; K97: 6.18 vs 3.50; R21: 1.03 vs 1.30) — direct evidence for the antigen-intrinsic hotspot idea (Design note 1). D1.3 binds a *non-overlapping* epitope (residues 18–24 and 116–129; zero residue overlap with the HyHELs), giving a second epitope plus built-in negative controls. Because we extract features from the apo antigen, **one apo-HEL trajectory** serves all three antibodies and their mean.

Starting structure. RCSB **1AKI** (1.5 Å X-ray, chain A, apo). Verified: the 1AKI sequence matches the benchmark WT identity at all 23 labeled positions, so labels map 1:1 onto 1–129 numbering (K96/K97/Y20 hotspots confirmed). Apo crystal chosen over antigen-stripped-from-complex to match the prediction scenario; HEL’s rigidity makes the choice low-risk.




Figure 1: Apo HEL starting structure (RCSB **1AKI**, 1.5 Å X-ray, chain A); cartoon coloured by secondary structure (helices purple, sheets yellow). This single apo trajectory serves all three antibody scans.

5 Simulated antigens

One section per antigen we simulate: what it is, the structural handle we use, and its alanine-scan labels. Each structure is coloured per residue by $\Delta\Delta G_{\text{bind}}$ on a diverging scale with **white pinned at 0**: red = positive (hotspot / likely epitope), blue = negative, grey = residue not scanned. One panel per antibody; a single global colour scale ($\pm\max|\Delta\Delta G_{\text{bind}}|$ over the whole benchmark, white at 0) is shared by every figure, so colours are directly comparable across antigens. Vertical colourbar at right. $\Delta\Delta G_{\text{bind}}$ is the change in binding free energy, K_d/K_i -derived; positive weakens binding.

5.1 Lysozyme (hen egg-white, HEL)

Hen egg-white lysozyme is a 129-residue glycoside hydrolase and the canonical model antigen of structural immunology: multiple high-resolution antibody complexes and alanine scans exist (HyHEL-10, HyHEL-63, D1.3). It is compact and conformationally rigid, cross-linked by four disulfides. We simulate the apo enzyme from **1AKI**. The labels here are the HyHEL-63 scan (1DQJ, Table 1) [7]; the v2 benchmark adds D1.3 (1VFB) [8] and HyHEL-10 (3HFM) [9] scans for the same antigen. These $\Delta\Delta G_{\text{bind}}$ values are drawn from the AB-Bind database [4].




Figure 2: HEL coloured per residue by alanine-scan $\Delta\Delta G_{\text{bind}}$ for its three antibodies (shared scale, white = 0, red = hotspot, blue < 0 ; grey = not scanned). Distinct but overlapping epitope cores across HyHEL-63, D1.3 and HyHEL-10.

Table 1: **Lysozyme** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; SPR, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310$ K; start structure: apo 1AKI). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	HyHEL-63 (1DQJ)	IgG1-kappa D1.3 Fv (1VFB)	HyHEL-10 Fv (3HFM)
D18	–	+0.30	–
N19	–	+0.30	–
Y20	+3.30	–	+4.87
R21	+1.30	–	+1.03
Y23	–	+0.40	–
S24	–	+0.80	–
W62	+0.80	–	–
W63	+1.30	–	–
L75	+1.50	–	–
T89	+0.80	–	–
N93	+0.60	–	–
K96	+6.10	–	+6.83
K97	+3.50	–	+6.18
S100	+0.80	–	–
D101	+1.30	–	–
K116	–	+0.70	–
T118	–	+0.80	–
D119	–	+1.00	–
V120	–	+0.90	–
Q121	–	+2.90	–
I124	–	+1.20	–
R125	–	+1.80	–
L129	–	+0.20	–

5.2 HPr

The histidine-containing phosphocarrier protein (HPr) is a small (~ 9 kDa, 85-residue) cytoplasmic protein of the bacterial phosphoenolpyruvate:sugar phosphotransferase system, shuttling a phosphoryl group between enzyme I and the sugar-specific EIIA components. The *E. coli* protein is a well-folded open-faced β -sandwich with *no cysteines*, making it an exceptionally clean dynamics test case. The Jel42 epitope clusters in the C-terminal half (Table 2) [10]; $\Delta\Delta G_{\text{bind}}$ values via AB-Bind [4]. We simulate the apo protein from 2JEL (chain P).




Figure 3: HPr coloured per residue by Jel42 alanine-scan $\Delta\Delta G_{\text{bind}}$ (white = 0, red = hotspot; grey = not scanned). A single dominant hotspot (E70).

Table 2: **HPr** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; competition, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310$ K; start structure: strip 2JEL chain P). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	Jel42 (2JEL)
T62	+0.00
E68	+0.41
E70	+2.75
H76	-0.41
E83	+0.00
E85	+0.00

5.3 VEGF-A

Vascular endothelial growth factor A is a secreted pro-angiogenic cytokine and a central drug target (bevacizumab, ranibizumab). Its receptor-binding domain is an antiparallel, disulfide-linked homodimer with a cystine-knot growth-factor fold; the two protomers are covalently joined, so the biological unit must be simulated as the dimer. The benchmark pairs two antibodies against the same receptor-binding face: fab-12 (1BJ1, the bevacizumab precursor) [11] and the affinity-matured

Y0317 (1CZ8, ranibizumab) [12] — a useful same-antigen/different-affinity comparison (Table 3); $\Delta\Delta G_{\text{bind}}$ values via AB-Bind [4]. We simulate the apo dimer from 1BJ1 (chains V, W).




Figure 4: VEGF-A homodimer coloured per residue by $\Delta\Delta G_{\text{bind}}$ for each antibody (shared scale, white = 0; grey = not scanned). Both protomers carry the same labelled residues; red marks the hotspots.

5.4 MT-SP1 (matrptase)

MT-SP1 / matrptase (ST14) is a type II transmembrane serine protease of the S1 family; its extracellular catalytic domain is over-expressed in epithelial cancers and is a validated oncology target. The benchmark uses two inhibitory Fabs, E2 (3BN9) [13] and S4 (3NPS) [14], scanning the protease domain in chymotrypsin numbering (Table 4; $\Delta\Delta G_{\text{bind}}$ via AB-Bind [4]). Simulation is staged but **pending**: the crystal structures lack two internal residues (149, 218) that must be repaired before MD (§11).

figures/mtsp1_structure.png
(not yet regenerated; see figures/FIGURES.md)

Figure 5: MT-SP1 protease domain (3NPS chain A). *Placeholder.*

Table 3: **VEGF** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; SPR, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310$ K; start structure: strip 1BJ1 chains V,W (dimer)). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	fab-12 (1BJ1)	Y0317 (1CZ8)
F17	+0.00	+0.00
Y21	+0.00	+0.00
Y45	+0.82	+1.94
K48	+0.41	+0.00
Q79	+0.00	+0.65
I80	+0.82	+0.95
M81	+3.69	+4.10
R82	+3.69	+0.82
I83	+3.69	+1.30
K84	+0.65	+1.36
H86	+0.00	+0.00
Q87	+0.00	+0.00
G88	+2.76	+2.67
Q89	+1.75	+1.06
H90	+0.00	+0.00
I91	+0.41	+1.06
G92	+3.69	+4.10
E93	+0.82	+1.15
M94	+1.42	+1.91

5.5 Bont/A1 (H_C receptor-binding domain)

Botulinum neurotoxin serotype A1 is among the most potent toxins known: a ~ 150 kDa di-chain protein comprising a zinc-metalloprotease light chain and a heavy chain. Its C-terminal heavy-chain domain (H_C , here residues 872–1295) is the receptor-binding domain — engaging neuronal gangliosides and the SV2 protein — and the principal protective antigen targeted by neutralizing antibodies. We simulate the isolated H_C domain (2NYY chain A); the catalytic light-chain zinc lies outside this fragment. Two neutralizing mAbs, CR1 (2NYY) and AR2 (2NZ9), provide the labels (Table 5) [15]; $\Delta\Delta G_{\text{bind}}$ via AB-Bind [4].

Table 4: **MT-SP1** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; enzymatic inhibition, K_i -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310 \text{ K}$; start structure: strip 3NPS chain A). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	E2 (3BN9)	S4 (3NPS)
Q23	–	+0.03
I26	–	+0.65
Q38	-0.42	–
I41	+0.00	–
I45	–	-0.34
D46	–	+0.34
D47	–	+1.08
R48	–	-1.07
F50	–	+0.22
R51	–	+0.14
Y52	–	+0.46
I60	+0.84	–
R82	–	-0.15
R87	-0.16	–
F89	–	+1.61
N90	–	+0.26
D91	–	+1.52
F92	–	+0.47
T93	–	+0.73
F94	+0.64	–
N95	+0.77	–
D96	+6.70	–
F97	+6.70	–
T98	+1.13	–
H138	–	+1.89
Q140	–	+0.30
Y141	–	+1.79
H143	+0.09	–
T144	–	+0.18
Q145	+0.13	–
Y146	+1.09	–
L147	–	+0.30
T150	+0.29	–
L153	+0.34	–
E163	–	+0.62
Q168	–	-0.06
E169	+0.37	+0.75
Q174	-0.03	–
Q175	+2.51	–
D214	–	+1.48
D217	+0.57	–
Q218	–	-0.04
R219	–	-0.08
K221	–	-0.10
R222	-0.09	–
K224	+0.79	–




Figure 6: Bont/A1 H_C coloured per residue by $\Delta\Delta G_{\text{bind}}$ for two neutralizing mAbs (shared scale, white = 0; grey = not scanned). Red marks the hotspots — the strongest reach $\Delta\Delta G_{\text{bind}} > 7$ (binding essentially abolished).


Table 5: **Bont/A1** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; KinExA, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310$ K; start structure: extract 2NYY Hc (res 872-1295)). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	CR1 (2NYY)	AR2 (2NZ9)
S902	-0.23	-0.12
K903	+0.29	+0.54
Q915	+0.52	+0.10
F917	+0.42	-0.05
N918	+0.89	+2.16
L919	+2.59	+2.28
E920	+2.84	+2.77
K923	-0.14	-0.30
F953	+4.06	+3.34
N954	+0.10	-0.15
S955	+0.00	-0.08
I956	-0.01	+0.07
K1056	-0.03	-0.01
D1058	+0.01	-0.03
R1061	+0.82	+0.29
D1062	+2.53	+2.34
T1063	+1.62	+2.37
H1064	+7.32	+7.42
R1294	+0.30	+0.39

5.6 IFN- γ receptor

The interferon- γ receptor α -chain (IFN- γ R1) ectodomain is the high-affinity ligand-binding subunit of the type II interferon receptor: two fibronectin type III domains forming an L-shaped β -sandwich

that engages the IFN- γ homodimer. The neutralizing antibody A6 binds this ectodomain and the published alanine scan localises a compact energetic hotspot on one face (Table 6) [16]; $\Delta\Delta G_{\text{bind}}$ values curated via SKEMPI 2.0 [6]. It is a clean, rigid, single-chain ~ 95 -residue Ig-like domain — an attractive small test antigen. We simulate the apo ectodomain stripped from 1JRH (chain I).



A6 (1JRH)

Figure 7: IFN- γ receptor coloured per residue by A6 alanine-scan $\Delta\Delta G_{\text{bind}}$ (shared scale, white = 0, red = hotspot; grey = not scanned). The hotspots cluster on a single β -sandwich face.

5.7 Tissue factor

Tissue factor (TF, coagulation factor III) is the cell-surface initiator of the extrinsic blood-coagulation cascade: its extracellular region is a two-domain fibronectin type III module that binds and allosterically activates factor VIIa. The inhibitory Fab 5G9 blocks the macromolecular substrate site, and its antigen-side alanine scan marks the functional epitope (Table 7) [17]; $\Delta\Delta G_{\text{bind}}$ values curated via SKEMPI 2.0 [6]. The ectodomain is a well-folded, single-chain ~ 200 -residue β -rich protein. We simulate the apo ectodomain stripped from 1AHW (chain C).

Table 6: **IFN-gamma receptor** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; SPR, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310\text{ K}$; start structure: strip 1JRH chain I). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	A6 (1JRH)
K47	+3.71
N48	+0.17
Y49	+3.53
G50	+4.45
V51	+1.68
K52	+3.39
N53	+4.30
S54	+0.38
E55	-0.57
N79	-0.20
W82	+4.43
R84	+0.39
K98	+0.31




Table 7: **Tissue factor** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; SPR, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310$ K; start structure: strip 1AHW chain C). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	5G9 (1AHW)
Y156	+3.81
Y157	-1.89
T167	-0.07
T170	+1.11
L176	+0.99
D178	-0.48
T197	+1.35
V198	-0.31
N199	+1.08

5.8 HCMV glycoprotein B

Glycoprotein B (gB) is the conserved class III fusion machine of human cytomegalovirus and a principal target of neutralizing antibodies. The ectodomain is a large, elongated trimer; the antibody 1G2 recognises antigenic domain 2 and its alanine scan identifies the contact hotspots (Table 8) [18]; $\Delta\Delta G_{\text{bind}}$ values curated via SKEMPI 2.0 [6]. This is our largest antigen (~ 566 modelled residues) and is biologically trimeric — we simulate the apo protomer as deposited in 5C6T (chain A), noting that the native oligomeric state may modulate the dynamics of the buried interfaces.




Figure 9: HCMV glycoprotein B protomer coloured per residue by 1G2 alanine-scan $\Delta\Delta G_{\text{bind}}$ (shared scale, white = 0, red = hotspot; grey = not scanned).

Table 8: **HCMV glycoprotein B** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; SPR, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310$ K; start structure: strip 5C6T chain A). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	1G2 (5C6T)
Y280	+3.21
N281	+0.28
T283	+0.34
N284	+0.36
N286	-0.37
F290	+2.00
E292	+1.09
F297	+2.76
F298	+0.41
I299	+0.89

5.9 Human growth hormone

Human growth hormone (hGH, somatotropin) is a 191-residue four-helix-bundle cytokine and the historical birthplace of alanine-scanning mutagenesis. Here the antibody-specific labels are the

MAb3 scan (Table 9) [5] [verify] (benchmark cites L. Jin’s 1994 thesis, Table 3.2; the published values are in [5]); the benchmark additionally carries a secondary, provisional profile of the total $\Delta\Delta G_{\text{bind}}$ aggregated across 21 anti-hGH MABs (not antibody-specific, held out of the primary labels). We simulate the apo hormone from 1HGU (chain A).




Figure 10: Human growth hormone coloured per residue by MAb3 alanine-scan $\Delta\Delta G_{\text{bind}}$ (shared scale, white = 0, red = hotspot; grey = not scanned). Only the four MAb3-scanned positions are coloured.

Table 9: **human growth hormone** alanine-scan $\Delta\Delta G_{\text{bind}}$ (kcal/mol; ELISA, K_d -derived via $\Delta\Delta G_{\text{bind}} = RT \ln(K^{\text{mut}}/K^{\text{wt}})$, $T = 310\text{ K}$; start structure: apo 1HGU). Positive $\Delta\Delta G_{\text{bind}}$ weakens binding.

Residue	MAb3 (1HGU)
R8	+2.16
R16	+3.68
D112	+1.84
D116	+2.73

6 A second label tier: binary critical-residue maps

The $\Delta\Delta G_{\text{bind}}$ antigens above are the gold standard — a graded, energetic importance per residue — but experimental alanine scans with full $\Delta\Delta G_{\text{bind}}$ tables are scarce. A much larger body of data exists at a coarser but still *energetic* resolution: **saturating shotgun alanine scans** read out by flow cytometry (Integral Molecular’s platform; Davidson & Doranz [19]), which classify each scanned residue as *critical* or *non-critical* for a given antibody. A residue is critical when its mutation measurably abolishes binding, so even this binary call is grounded in binding energy, not mere proximity — it is the same quantity as $\Delta\Delta G_{\text{bind}}$, thresholded. We therefore treat these as a second label tier: **Tier 1**, graded $\Delta\Delta G_{\text{bind}}$ (§5); **Tier 2**, binary critical/non-critical.

Source and the proximity filter. We mine these from IEDB (query-api.iedb.org) [1], keeping only conformational (discontinuous) epitopes defined by *binding* assays and explicitly discarding every structure/contact-derived record — the proximity labels this project argues against (recall that $> 99\%$ of structure-linked IEDB epitopes are crystal/cryo-EM contact, §1). Within the binding-derived set we take the flow-cytometry (shotgun-mutagenesis) subset, whose curated residue list *is* the set of critical residues. This yields 3,213 antibody–antigen epitope maps over 186 antigens. Table 10 lists the simulatable, not-yet-included candidates; note the depth of antibody sampling per antigen (e.g. dengue envelope: 781 mapped antibodies across 30 studies), which lets a Tier-2 antigen carry a *consensus* epitope (residues critical across many antibodies) as well as antibody-specific maps.

Table 10: **Tier-2 (binary) label candidates** from IEDB shotgun-mutagenesis alanine scans (flow-cytometry critical-residue mapping; Integral Molecular / Doranz). Labels are per-residue *critical vs non-critical*, energetically grounded (a critical residue is one whose mutation measurably weakens antibody binding) but coarser than $\Delta\Delta G_{\text{bind}}$. “Antibodies” = mapped antibody footprints (one antibody $\rightarrow$ one conformational epitope); “Studies” = distinct source papers; “Crit. res.” = union of critical residues across all antibodies. Structures provisional ([v] = numbering to be verified before use).

Antigen	Organism	Antibodies	Studies	Crit. res.	Structural handle
Dengue E (envelope)	Dengue virus	781	30	632	sE ectodomain, e.g. 1OKE [v]
Chikungunya E1/E2	Alphavirus chikungunya	128	6	98	ectodomain, e.g. 3N42 [v]
RVFV Gn	Rift Valley fever virus	63	8	210	Gn ectodomain, e.g. 6F9F [v]
Influenza HA	Influenza A virus	81	22	296	HA ectodomain (trimeric) [v]
GRO- α / CXCL1	Homo sapiens (human)	30	1	23	chemokine, e.g. 1MGS [v]
CD4	Homo sapiens (human)	1	1	18	D1–D2, e.g. 1CDH [v]
CD40L	Homo sapiens (human)	1	1	13	TNF domain, e.g. 1ALY [v]
IL-2R α	Homo sapiens (human)	2	1	13	ectodomain, e.g. 1Z92 [v]
IL-7R α	Homo sapiens (human)	2	1	32	ectodomain, e.g. 3DI2 [v]
HLA class I	Homo sapiens (human)	15	5	38	peptide-MHC, many [v]

Numbering verification. The shotgun residues are reported in UniProt numbering and carry their wild-type identity, but aggregating by UniProt accession mixes proteins, numbering conventions, and (for viruses) serotypes/subtypes. We therefore map every residue onto a single reference structure and *discard any residue whose wild-type identity does not match* — exactly the wild-type check applied to the $\Delta\Delta G_{\text{bind}}$ antigens. This is strict: of dengue’s raw critical positions only those matching DENV2 envelope survive, and likewise only H3 residues survive for influenza. What remains is clean and biologically coherent — the top dengue residue (fusion loop, E101) is called critical by 102 of 632 antibodies.

Representation. We render each staged antigen two ways (Figs. 11, 12): **binary** (red = critical for at least one antibody, grey = not) and **consensus frequency** (white $\rightarrow$ red by the number of antibodies for which the residue is critical). The frequency view is itself an energetic-importance signal — it ranks residues by immunodominance across hundreds of antibodies — and reduces to the binary view under any threshold, so we keep both and decide later. Tier-2 antigens enter the benchmark as a separate (binary/graded) label stream, never merged into the $\Delta\Delta G_{\text{bind}}$ tables; in the leakage-controlled evaluation they are held out by antigen.




Figure 11: **Dengue envelope (DENV2, 10KE sE)**: 96 wild-type-verified critical residues from 632 antibodies. Binary (left) and antibody-frequency consensus (right). The deepest red is the conserved fusion loop (E101, critical for 102 antibodies); a second focus is the EDIII lateral ridge — both textbook dengue neutralization sites.




Figure 12: **Influenza hemagglutinin (H3 X-31, 1HGG protomer)**: 84 wild-type-verified critical residues (HA1 head + HA2 stem) from 79 antibodies, binary and antibody-frequency consensus. Critical residues concentrate on the HA1 globular head (the classic antigenic sites) with a sparser stem contribution. Native HA is an obligate trimer; one protomer is shown.

7 Static baselines: exposure and mobility

Before adding any dynamics, we ask how far purely *static*, single-crystal descriptors already get — both to set the bar the dynamics features must clear and to make the proximity critique of §1 quantitative on our own labels. We test the two classic static proxies for “epitope-ness”: surface **exposure** (this section) and **mobility** (crystallographic B-factors, below). The exposure descriptor is per-residue **relative solvent-accessible surface area** (rel. SASA): Shrake–Rupley SASA on the input crystal (antigen chain only), normalised by the residue-type maximum [20]. This is a single-frame calculation on the input crystal structures — the pre-trajectory baseline; once the production runs land it is replaced by trajectory-averaged SASA at the same residue keys (and joined by RMSF, non-bonded energetics, and hydration).

Which structures (a confound shared by both baselines). Both descriptors are read from the deposited crystal of each antigen, and for *six of the eight* the crystal is the antibody–antigen **complex** — only lysozyme (1AKI) and hGH (1HGU) are true apo structures. This biases the two features differently. For SASA we compute on the antigen chain *alone* (antibody atoms deleted), so the epitope surface is not buried in the calculation — but the coordinates are the *bound* conformation, so these are holo-state SASA values, equal to apo only absent induced fit. For B-factors the bias is more direct — the bound antibody damps exactly the epitope residues’ B-factors — so we re-test on genuine apo crystals in the mobility section below. The SASA panel remains single-crystal and reads as a motivating bound, not a clean measurement.

Setup. We pool every residue of the eight $\Delta\Delta G_{\text{bind}}$ antigens (1,790 residues). A residue is *important* if its alanine-scan $\Delta\Delta G_{\text{bind}} \geq 1$ kcal/mol for any antibody (aggregated as the max over antibodies). An exposure-only predictor labels a residue an epitope if rel. SASA ≥ 0.25 (the standard “exposed” cutoff) — the static analogue of a proximity label. We then read off precision and recall against the experimental importance labels.




Figure 13: **Surface exposure as an epitope label.** Every residue of the eight $\Delta\Delta G_{\text{bind}}$ antigens, rel. SASA (x , antigen chain alone) vs. alanine-scan $\Delta\Delta G_{\text{bind}}$ (y); red = scanned and important ($\Delta\Delta G_{\text{bind}} \geq 1$), blue = scanned but weak ($\Delta\Delta G_{\text{bind}} < 1$, incl. binding-improving), grey = not alanine-scanned (no data, plotted at 0); the shaded band is what an exposure cutoff (rel. SASA ≥ 0.25) would call epitope. Exposure flags 925 of 1,790 residues as epitope but only 31 are truly important (precision $\approx 3\%$), while still missing the buried-at-interface hotspots to the left of the line (recall $\approx 63\%$). Single-frame SASA on the input crystals (6/8 complex-derived, see text); `analysis/compute_sasa.py` + `plot_feature_vs_label.py`.

Result. Exposure behaves exactly as the thesis predicts. It has *moderate recall* (it knows roughly where the surface is: 63% of importance hotspots are exposed) but *near-zero precision* ($\approx 3\%$): rel. SASA ≥ 0.25 sweeps in almost half the protein, overwhelmingly residues that contribute no binding energy. Per antigen the precision is 0–14% (highest for the clean single-domain antigens 1AKI, 1BJ1, 1JRH). And *among the scanned interface residues*, exposure does not even rank graded importance: the pooled Spearman correlation between rel. SASA and $\Delta\Delta G_{\text{bind}}$ is -0.16 — once a residue is on the surface, how exposed it is says nothing about how much binding energy it carries. That gap — high recall, no precision, no within-interface ranking — is precisely what the dynamic and energetic features are meant to fill, and every feature we add is scored against this baseline.

Caveats. (i) Single-frame SASA on the input crystals (6/8 complex-derived, so bound conformation; antibody removed for the calc); trajectory-averaging on the apo simulation is the real control.

(ii) 5C6T (gB) is anomalous here — it is scored on the unrepaired structure (§11) *and* as an isolated protomer of a native trimer, so residues buried in the trimer interface read as exposed; 2JEL (HPr) carries only one hotspot. (iii) Residues never scanned are treated as non-epitope, which alanine scans largely justify (they target the known epitope) but which inflates the negative set.

Static mobility (crystallographic B-factors) — the dynamics teaser. A second, long-standing idea links antigenicity to *flexibility*: mobile surface regions are more antigenic [2]. The static, experimental proxy for residue mobility is the crystallographic B-factor, and it is the *static analogue of the RMSF* we will extract from the apo trajectories — so a B-factor signal would foreshadow the dynamics payoff. We test it the same way: per-residue C α B-factor, z-normalised per structure (raw B-factors are incomparable across crystals — resolution, refinement, overall scaling), against $\Delta\Delta G_{\text{bind}}$ (Fig. 14). Because six of eight benchmark crystals are antibody complexes whose bound antibody damps the epitope residues’ B-factors, we run this test on **genuine apo structures**: HPr 10PD, VEGF 4KZN, tissue factor 2HFT, BoNT/A1 H_C 3FU0, gB 5CXF, plus the already-apo lysozyme 1AKI and hGH 1HGU (IFN- γ receptor dropped — no ligand-free structure exists; `analysis/recompute_apo_baselines.py`, all verified to cover the scanned residues). The result runs *against* the naive expectation, if weakly: the energetic hotspots are not the most mobile residues but trend toward the *rigid* side — pooled scanned Spearman $\rho = -0.23$ (-0.29 on the original complex-bound set, so the antibody-damping confound is real but minor). The trend is *not* robust: per antigen it swings from $\rho = +0.28$ (VEGF, mobile hotspots) to -0.70 (tissue factor). The honest statement is that crystal mobility gives *no* reliable flexible-hotspot signal, at most a weak rigidity tendency — which still cuts against the naive flexibility–antigenicity reading [2], whose correlation concerns where antibodies *bind* (surface loops), not which residue carries the binding *energy*, the proximity-vs-energetics distinction of §1. A clean answer needs the lattice- and antibody-free measurement that apo-MD RMSF provides.




Figure 14: **Crystallographic mobility as an epitope label (apo structures).** Per-residue C α B-factor (z-scored per structure; right = more mobile) on genuine apo (antibody-free) crystals vs. alanine-scan $\Delta\Delta G_{\text{bind}}$, colours as in Fig. 13. The important (red) residues lean toward *low* B-factor (rigid) rather than the high-mobility band, but only weakly (pooled Spearman -0.23 ; -0.29 on the complex-bound set) and inconsistently across antigens. `analysis/recompute_apo_baselines.py` + `plot_feature_vs_label.py`.

Flexibility versus energetic isolation. That the hotspots lean rigid rather than mobile — weakly, but with no hint of the opposite — is less a contradiction of the flexibility hypothesis than a pointer to a different variable. Westhof’s mobility [2] and our B-factor test both ask about *motion*; but the property that makes a surface patch easy for an antibody to engage may instead be *energetic and dynamic isolation* — the thesis of Serapian and Colombo [3], whose MD-based energy decomposition (a matrix of local inter-residue coupling energies) identifies antibody epitopes as surface regions *weakly coupled* to the rest of the protein: substructures that make few stabilising nonbonded contacts with the core, sit near energetic frustration, and behave as semi-autonomous dynamic units. Such a patch can be grabbed without fighting the protein’s own stability network — and if it is also locally rigid (pre-organised), the antibody pays little conformational-entropy penalty to hold it: *decoupled enough to grab, ordered enough to keep*. On this view the predictive quantity is not raw mobility but a residue’s *coupling* to its environment — intra-protein nonbonded interaction energy and the correlation structure of its motions. Those are exactly the features the apo-MD program targets (intra-protein energetics, correlated motion, FRESEAN mode participation), and they are blind spots of both static baselines here, which see exposure and motion but not coupling.

Where this leaves us. Neither static proxy recovers the energetic hotspots from a single crystal: exposure has no precision ($\sim 3\%$), and apo mobility gives no reliable signal (weak, inconsistent $\rho \approx -0.23$). That is precisely the opening for apo-MD — RMSF is a cleaner, lattice- and antibody-free mobility measure, and the trajectory additionally exposes the correlated-motion, mode-participation (FRESEAN), and hydration (3D-2PT) features that no single structure carries. B-factors are a reasonable first-pass proxy — the standard one — but even re-run on apo crystals they conflate intrinsic flexibility with resolution, refinement, and packing; the per-structure z-score removes only the overall scale. The question the rest of this work asks is whether dynamics *done properly* recovers signal where these static proxies fail.

8 Dynamics, first feature: apo-MD RMSF

The B-factor test (§7) was the static teaser; this is the measurement it was standing in for. With the apo production runs landed, we replace crystallographic mobility with its lattice- and antibody-free dynamic analogue: per-residue **root-mean-square fluctuation** (RMSF) from the apo trajectories. If the weak rigidity lean of the B-factors ($\rho = -0.23$) was a crystallographic artefact — resolution, refinement, packing — proper dynamics should wash it out; if it was real, RMSF should sharpen it. Either way this is the first feature of the dynamics program, and the one most directly comparable to a static baseline.

Setup. Per-residue $C\alpha$ RMSF over the equilibrated portion of each 100 ns apo trajectory (first 10 ns discarded; least-squares fit to the reference, so global tumbling is removed before fluctuations are measured), z-normalised per structure to match the B-factor convention (`analysis/rmsf_dynamics_baseline`, GROMACS 2023.2 on Gemini). Labels are kept at their native granularity — one column per (antigen, antibody) *epitope*, never averaged over antibodies, so lysozyme’s three antibodies (HyHEL-63, HyHEL-10, and the opposite-face D1.3) and VEGF’s two are scored separately. Correlation is reported per epitope, pooled over all scanned residues, and **leave-one-antigen-out** (LOAO): the honest generalisation number, since residues within an antigen share a fold and a trajectory and our effective N is the number of antigens, not residues (§2). The split is by antigen, not antibody — lysozyme’s three antibodies share one apo trajectory, so an antibody-level split would leak features across train and test.

Which systems. Six of the finished apo runs align cleanly to the labels (1AKI, 2JEL, 1BJ1, 1JRH, 1AHW, 1HGU); the MD structure is numbered differently from the antibody-complex crystals the scans come from, so alignment is *verified by residue identity* — each scanned wild-type residue must match the amino acid actually present at that position in the simulated structure, not assumed from an offset. All six systems align at zero offset. Two graded-scan antigens are deliberately *not* in this set: **MT-SP1** (3BN9/3NPS) has $\Delta\Delta G_{\text{bind}}$ labels but no apo-MD trajectory yet, and **Dengue envelope** (10KE) has an apo trajectory but Tier-2 *binary* labels (§6), which are scored on the binary track rather than joined to the graded set here. The influenza HA run (1HGG, also Tier-2) is the largest system and had not finished at the time of this iteration. That leaves 9 antibody-epitopes over 6 antigens, 97 scanned residues.

Real apo-MD flexibility (RMSF) mirrors the B-factor baseline: no signal beyond exposure
6 antigens · 9 antibody-epitopes · 97 scanned residues · pooled $\rho = -0.12$ ($p = 0.24$) · LOAO mean $\rho = -0.15$ [MT-SP1/10KE excluded; numbering unaligned]




Figure 15: Apo-MD flexibility as an epitope label. *Left:* Spearman ρ between per-residue apo RMSF and alanine-scan $\Delta\Delta G_{\text{bind}}$ for each antibody-epitope (coloured by antigen; dashed line is the pooled value); the correlations scatter around zero with a slight negative lean. *Right:* all 97 scanned residues, RMSF (z ; right = more mobile) vs. $\Delta\Delta G_{\text{bind}}$, colours as in Fig. 14 — the important (red) residues are spread across the mobility range, not concentrated on the mobile side. An $\text{RMSF} \geq 0$ cutoff (the shaded band) flags 55 of 97 scanned residues and catches 25 of 48 hotspots — no better than a coin toss on this set. `analysis/rmsf_dynamics_baseline.py`. MT-SP1 (no apo MD) and Dengue E (10KE, Tier-2) are not in this graded set; HA (1HGG) pending.

Result. Proper dynamics does not rescue the flexibility hypothesis — it confirms the B-factor verdict on cleaner data. The pooled scanned Spearman between RMSF and $\Delta\Delta G_{\text{bind}}$ is $\rho = -0.12$ ($p = 0.24$, $n = 97$), and the leave-one-antigen-out mean is $\rho = -0.15$ (over the five antigens with ≥ 5 scanned residues) — statistically indistinguishable from zero, and the same weak-rigid *sign* the apo B-factors gave (-0.23). Per epitope the correlations run from $\rho = -0.80$ (hGH, lysozyme HyHEL-10, both $n = 4$) through -0.66 (lysozyme D1.3) to $+0.29$ (IFN- γ receptor); the sign is not even consistent across the three antibodies of a single antigen. As a classifier the picture is the same: thresholding at $\text{RMSF} \geq 0$ recovers roughly half the hotspots at roughly half the scanned residues — the diagonal of no information. Whatever distinguishes a binding-energy hotspot from its neighbours, it is not how much that residue moves.

Reading. This is a negative result on the feature, but a positive one for the thesis of §7: neither static nor dynamic *mobility* — B-factor or RMSF — carries the hotspot signal, which is exactly what the energetic-isolation view predicts. Motion is the wrong observable; what should matter is a residue’s *coupling* to the rest of the protein, not its amplitude of fluctuation [3]. RMSF measures the amplitude and finds nothing, cleanly and on real trajectories. The same runs that produced these fluctuations carry the observables that the coupling view actually points at — intra-protein nonbonded interaction energy, the correlation structure of the motions, and mode participation — and those are the next features, scored against this same baseline.

Caveats. (i) Single 100 ns trajectory per antigen; RMSF of the most flexible loops is not fully converged at this length, though the hotspot residues (ordered, low-RMSF) are the stable part of the estimate. (ii) Two epitopes rest on only $n = 4$ scanned residues (hGH, lysozyme HyHEL-10), so their individual ρ values are indicative, not reliable; the pooled and LOAO numbers are the ones to trust. (iii) MT-SP1 (no apo MD yet) and the Tier-2 systems Dengue E and HA (1HGG run pending) are absent from this graded set — the pooled value will be re-computed as systems join. (iv) As with the static baselines, unscanned residues are treated as non-epitope.

9 From amplitude to coupling: energetic and collective features


The RMSF result (§8) closed by naming the observables the energetic-isolation view actually points at — not how much a residue moves, but how it is *coupled* to the rest of the protein: its intra-protein interaction energy, the correlation structure of the motions, and its participation in the collective modes [3]. The same apo runs that gave the fluctuations carry all three. This section scores them, together with a geometric hydration count, against the same graded $\Delta\Delta G_{\text{bind}}$ baseline — and reaches the first feature that is not simply zero.

Setup. Every feature is per-residue, from the equilibrated portion of each 100 ns apo trajectory (first 10 ns discarded), z-normalised per structure, on the six antigens that align to the graded scans at zero offset (§8). **Intra-protein interaction energy** is the residue’s summed nonbonded (Coulomb + Lennard-Jones) energy with every other residue of its *own* fold, $E_i^{\text{intra}} = \sum_{j \neq i} (E_{ij}^{\text{elec}} + E_{ij}^{\text{vdW}})$, evaluated frame-by-frame with gRINN (its own bundled GROMACS, `amber99sb-ildn`; the antibody is absent, so this is a property of the apo antigen, not an interface energy). **DCCM coupling** is the mean absolute dynamic cross-correlation of a residue’s $\text{C}\alpha$ displacement with all others. **PC1 participation** is the residue’s weight in the slowest mode of the $\text{C}\alpha$ position covariance — a placeholder for the collective picture, and one I read with the caveat below. **Hydrogen bonds** are counted geometrically ($D \cdots A \leq 3.5 \text{ \AA}$, $D\text{--}H \cdots A \geq 150^\circ$) with explicit `amber99sb-ildn` donor/acceptor names, split into protein–protein (the internal network) and protein–water (hydration). Scoring is leave-one-antigen-out throughout, as in §2.

Score on the surface, and read the direction. Two changes to the scoring earn their keep here. First, the honest question on this label is not whether a feature separates the whole protein — that is dominated by the buried/exposed split, which every feature reports and which is not what we are after. Restricting to *surface* residues (trajectory SASA above the per-antigen median) removes that confound, and it is where exposure finally shows its hand: within the surface population, trajectory SASA separates hotspots from their neighbours at AUROC = 0.51 — a coin toss. Raw accessibility only ever found the outside of the protein, exactly as the energetic-isolation thesis (§7) predicts it should. Second, the *sign* of the separation is the result, so I report the directional AUROC = $P(\text{hotspot} > \text{rest})$, not its folded magnitude (Fig. 16).

Which apo-antigen features distinguish antibody hotspots — and in which direction

Surface residues only (exposure confound removed) · 6 Tier-1 antigens · hotspot = alanine-scan $\Delta\Delta G \geq 1$ kcal/mol · bar length = separation strength · bold+outlined = new H-bond features



H-bonds to water (0.34) join PC1 participation (0.32) and backbone flexibility (0.33) as the strongest LOW-direction features: among equally exposed residues, hotspots are less hydrated, sit stiller, and are decoupled from the soft mode — yet more internally bonded (interaction energy, DCCM, internal H-bonds). All weak at $N=6$.

Figure 16: Which apo features separate hotspots from other surface residues, and in which direction. Surface residues only; bars left of centre = hotspots score *lower*, right = *higher*; length is the separation. Hotspots are less hydrated (fewer water H-bonds), sit at the quiet nodes of the slowest collective mode (low PC1 participation) and have stiffer backbones, yet are more internally bonded (intra-protein H-bonds, DCCM coupling, interaction energy). Every effect is weak ($N = 6$ antigens); the direction, not the magnitude, is the content. `analysis/feature_direction.py`.

A picture, not a predictor. Read as directions, the features agree on a coherent residue. The hotspot is *less* hydrated than its surface neighbours (protein–water H-bonds down), sits at a *node* of the slowest collective motion rather than swinging with it (PC1 participation down), and carries a *stiffer* backbone — while at the same time being held *more* tightly by its own fold (intra-protein H-bonds, DCCM coupling and interaction energy all up). These are not four independent readings: being well-bonded internally is what makes a residue rigid, and a rigid residue is what sits at the node of a soft mode. It is the same anchor the coupling view asks for, seen from four sides, and it lines up with the entropic argument of §7 — a pre-organised, rigid, desolvation-ready residue pays little on binding and can carry more of the $\Delta\Delta G_{\text{bind}}$. What it is *not* is a predictor: on this surface-restricted binary test the features cluster in a weak 0.5–0.65 band, and at $N = 6$ the ordering within that band is noise.

The label is graded — so grade it. Dichotomising a continuous $\Delta\Delta G_{\text{bind}}$ at 1 kcal/mol^{−1} throws away most of what the alanine scans measured. Correlating each feature against the graded $\Delta\Delta G_{\text{bind}}$ instead (per-antigen Spearman, the leakage-controlled number) reorders the field, and the reordering is the informative part. Interaction energy, which *led* the binary surface test ($\text{AUROC}_{\text{vdW}} = 0.65$), goes flat against the gradient ($\rho \approx +0.08$): it flags the handful of extreme

hotspots but does not track intermediate binding energy. Protein–water H-bonds do track it, and become the strongest single feature ($\rho = -0.48$, per antigen; pooled -0.40 , $p = 0.01$) — the more a surface residue hydrogen-bonds to solvent, the lower its $\Delta\Delta G_{\text{bind}}$. And it is *hydration*, not chemistry: static residue-type hydrophobicity (Kyte–Doolittle) is flat ($\rho \approx 0.00$), and the water count correlates with SASA at only $+0.19$, so this is context-dependent desolvation measured in the trajectory, not a property the sequence could have given (Fig. 17).




Figure 17: **Grading the label beats thresholding it.** Per-antigen Spearman ρ of each feature against continuous $\Delta\Delta G_{\text{bind}}$ (surface residues; pooled ρ overlaid as points). Hydration leads; interaction energy — the binary winner — is flat on the gradient. *Right:* the water-H-bond vs. $\Delta\Delta G_{\text{bind}}$ relation, coloured by antigen. `analysis/graded_ddg.py`.

One combination that does better — and does not overfit. The two strongest low-direction features measure genuinely different things: hydration is desolvation, PC1 participation is collective rigidity, and they are only weakly related residue-to-residue. Their sign-aligned, *equal-weight* sum $z(\text{HB}_{\text{water}}) + z(P^{(1)})$ — no fitted coefficients — tracks the graded $\Delta\Delta G_{\text{bind}}$ at per-antigen $\rho = -0.60$, above hydration alone (-0.48) and PC1 alone (-0.31), and negative in all four antigens with enough scanned surface residues to score (Fig. 18). The fixed weighting is the point: the same features under a *fitted* leave-one-antigen-out logistic go *below* chance, weights learned on five antigens pointing the wrong way on the sixth — the signature of overfitting at $N = 6$. Averaging two real, complementary signals with no free parameters is the honest version of “a combination that does better,” and adding a third feature degrades it, which says the information lives specifically in the desolvation and rigidity axes.




Figure 18: **Hydration + collective motion.** The equal-weight, unfitted collective variable $z(\text{HB}_{\text{water}}) + z(P^{(1)})$ vs. $\Delta\Delta G_{\text{bind}}$, centred within antigen so the pooled slope matches the per-antigen effect (per-antigen $\rho = -0.60$); the pooled uncentred cloud inverts by Simpson's paradox and is not the honest view. *Right:* within-antigen ρ for the combination against its two components. `analysis/combined_cv.py`.

Reading. This is the first feature line that is not flat, and it says something specific: the hotspot signal on an energy label is not motion and not exposure but *desolvation coupled to rigidity* — a residue that has given up its water and been locked into its fold. That is the energetic-isolation picture made concrete at the residue level, and it is worth stating where it sits relative to the patch-level version (§7): Colombo's argument predicts *where* on the surface an antibody can grab — a weakly-coupled, protrudable patch — while this result says *which residues within reach* carry the binding energy — the well-bonded, pre-organised anchors. Location and anchor, not a contradiction, but the same word “coupling” pointing at two different scales, and the manuscript should keep them apart.

Caveats. (i) $N = 6$ antigens, and the per-antigen correlations rest on four of them (tissue factor, VEGF, lysozyme, IFN- γ receptor); the others carry too few scanned surface residues for a within-antigen ρ . Every ranking here is a hypothesis, not an estimate. (ii) PC1 participation is Cartesian-covariance PCA, which partly re-measures amplitude (it correlates with RMSF at +0.52); the “hotspots at the nodes” claim only becomes mechanistic if it survives the frequency-selective FRESEAN modes now running, which is the correct operator for the soft collective motion and the reason those runs exist. (iii) The interaction energies are the short-range decomposable part; the long-range reciprocal-space Ewald term is not per-pair separable and is not in E^{elec} . (iv) The bottleneck is the label set, not the model — no fitted combination beat the best hand-chosen fixed one — which is why the pseudo-label scale-up matters more than any further feature.

10 Methods — apo-HEL run #1

Frozen, signed-off stack in `../docs/METHODS.md` (2026-06-18). Purpose of run #1: *validate the prep→equil→production pipeline on Gemini and extract a first RMSF / flexibility feature* (the Westhof re-test). House default stack signed off as the starting point.

Stack. AMBER99SB-ILDN / TIP3P; cubic box, 1.2 nm clearance; neutralize + 0.15 M NaCl; leapfrog, 2 fs, LINCS on h-bonds; PME ($r_{\text{coul}} = r_{\text{vdw}} = 1.0$ nm, potential-shift, DispCorr); V-rescale at 300 K (Protein / Non-Protein); EM $\rightarrow$ 100 ps NVT $\rightarrow$ 100 ps NPT (both position-restrained) $\rightarrow$ **100 ns** production (single trajectory, for RMSF statistics). Random seeds (**gen_seed** = **ld_seed** = -1): bitwise reproducibility is not the target — the workflow reproduces from git-tracked inputs and saved derived outputs, and random seeds give genuinely independent replicas. Frozen per-run .mdp files: md/1AKI/apo/configs/ (run dirs are organized md/<PDB>/<state>/).

Deliberate deviations from the house scripts. (i) **C-rescale** barostat for NPT equilibration (Parrinello-Rahman only in production), since PR can ring when started far from equilibrium; (ii) **cubic** box (the old house scripts disagreed cubic/dodecahedron) — chosen for interpretability and uniform across all systems, at a modest extra-solvent cost vs. dodecahedron.

HEL-specific checkpoints (not glossed over). pdb2gmx must form HEL’s **4 disulfides** (6–127, 30–115, 64–80, 76–94) — verify the log reports 4 SS bonds. His15 protonation: record the assigned state. Crystal waters (78 HOH) removed; fresh TIP3P added. Read every grompp warning rather than letting **-maxwarn** hide it.

Caveats / flags. TIP3P over-mobilizes water and 2 fs+LINCS suppress H dynamics — fine for relative RMSF and pipeline validation, *not* for the dynamics-grade features. 100 ns is a single trajectory for statistics; replicas are deferred (random seeds make adding independent ones trivial).

Deferred: the distinctive run. FRESEAN mode participation and 3D-2PT hydration thermodynamics require a *separate* protocol — high-frequency velocity output, constraint removal / sub-fs dt, TIP4P/2005 — with settings pulled from the FRESEAN and energy-transfer manuscripts. This is the run that powers the mode-based and hydration feature classes.

11 Structure completeness and the repair strategy

Several crystal structures are incomplete — disordered termini, internal loops, or whole segments that the experiment never resolved (PDB REMARK 465). Apo-MD demands a chemically complete chain, so each gap forces a choice. We make that choice explicitly per structure rather than feeding everything through one tool, because **the right repair scales with gap length**, and the wrong tool on a long gap produces a confident-looking but meaningless model.

What a “fixer” can and cannot supply. The instinct to “just grab the FASTA for the missing residues” conflates two different problems. The deposited sequence (**SEQRES**, the same residues a fixer recovers automatically by diffing **SEQRES** against the resolved **ATOM** records) supplies the *identity* of the missing residues — but identity is the easy half. The hard half is *geometry*: where the missing backbone atoms sit in three dimensions, about which the sequence says nothing. **pdbfixer** (OpenMM) builds missing heavy atoms onto resolved residues from templates — valid and routine — and threads missing residues in with default internal coordinates followed by a short minimization of only the new atoms. That minimization removes clashes; it performs *no* conformational sampling. The result is therefore reliable for a short surface gap (apo equilibration relaxes a two- or three-residue loop anyway) and arbitrary for anything longer, where real loop modelling (MODELLER, Rosetta) or a full predicted model (AlphaFold) is required. Disordered *termini* are not modelled at all: they bear no labels and would only add unrestrained flailing, so we cap them.

Method used: AlphaFold-loop grafting with restrained closure. Rather than let a fixer build loops blind, we graft the missing internal residues from an AlphaFold model and use the crystal for everything resolved. For each antigen we (i) sequence-align the crystal chain to its AlphaFold model to fix the numbering offset, (ii) measure the C α -RMSD between model and crystal over all matched residues — the “is the prediction trustworthy here” test — and (iii) for each internal gap, locally superpose the model on the flanking resolved residues and splice in its loop at crystal numbering. `pdbfixer` then completes any partial side chains, and a short OpenMM minimization with all crystal atoms position-restrained (only the loop free, plus an explicit harmonic across the spliced junction) closes the peptide bonds. Disordered termini are capped, not modelled. Scripts: `analysis/repair_structure.py`, `finalize_structure.py`, `close_loops.py`; host-side env `env/repair-environment.yml` (Python 3.11, OpenMM 8.5, `pdbfixer`, run locally, not on Gemini).

Antigen (PDB, chain)	Unresolved segments (REMARK 465)	Repair outcome
hGH (1HGU A)	F1; K38–E39 (internal); F191	AF-graft (AFDB P01241); done
Tissue factor (1AHW C)	S1–T3; V83–G90 (8-res); Q212–E219	AF-graft (AFDB P13726); done
HCMV gB (5C6T A)	63–87 (25); 115–120; 219–220; 237–239; 437–468 (32)	not in AFDB → AF3 prediction (pending)

For tissue factor the AlphaFold model agrees closely with the crystal (global C α -RMSD 1.12 Å over 200 residues; the V83–G90 loop fits its flanks at 1.29 Å), so the 8-residue graft is well supported. For hGH only the 2-residue K38–E39 turn is grafted (local fit 1.81 Å); the larger global RMSD (3.36 Å, hinge flexibility of the four-helix bundle) is immaterial for so short a splice. Both repaired chains validate as continuous, with the expected disulfides (hGH 53–165, 182–189; TF 49–57, 186–209) and no steric clashes. HCMV gB is a viral protein *absent from AlphaFold DB*, and its largest gap (32 residues) is true *de novo* prediction, not loop repair; it requires a fresh model. We generate one with **AlphaFold3** on Gemini (the lab’s `alphafold_3.0.1` Singularity container and installed weights/databases; `md/5C6T/apo/af3.sbatch`, mirroring the episcaf pipeline), then graft its loops the same way. The AF3 model’s agreement with the crystal over the resolved residues is the same trust check applied above; if it disagrees we simulate the AF3 ectodomain directly.

Consequences for the run queue. Eight of the nine non-pristine/pristine antigens are now ready to prep on Gemini (`git pull` then `sbatch`): pristine chains HEL (1AKI), HPr (2JEL), VEGF (1BJ1), IFN- γ R (1JRH), Dengue E (10KE), HA (1HGG); plus the repaired hGH (1HGU) and tissue factor (1AHW). Outstanding: HCMV gB (5C6T, awaiting an AF3 model) and the previously deferred MT-SP1 (§5, two internal residues, same AF-graft recipe) and Bont/A1.

12 Run log

Exact-command reproducibility ledger. Append a row per stage; pin the GROMACS version actually used on Gemini, seeds, and Slurm job id.

Date	Stage	Command	GROMACS ver.	Job / output
2026-06-18	structure	<code>curl .../1AKI.pdb</code>	—	<code>structures/1AKI.pdb</code>
<i>pending</i>	prep	<code>sbatch md/1AKI/apo/prep.sh</code>	2023.2-dev	<code>out/npt.gro/cpt</code>
<i>pending</i>	production	<code>sbatch md/1AKI/apo/md.sh</code>	2023.2-dev	<code>out/md.xtc</code>

Status (2026-06-18). Project scaffolded; benchmark audited; first system (HEL/1AKI) selected and label mapping verified 1:1; apo-HEL run #1 staged and signed off. Gemini environment confirmed: module `Gromacs` (capital G), GROMACS 2023.2-dev (avx2_256), partition `gpu-a100` (4-day limit, A100 fastest); repo cloned to `/scratch/bneff/bcell_epitope`. Production walltime set to 48 h ceiling. **Next:** `git pull` on Gemini, then `sbatch prep.sh` → `sbatch md.sh`; record Slurm job ids and ns/day. Production length 100 ns (single trajectory). Open: replicas (deferred — random seeds make them trivial); scheduling the FRESEAN/3D-2PT run.

Run-script fixes (bringing up run #1 on Gemini). Three issues surfaced and were fixed while getting prep to run; recorded for reproducibility:

- **Module name.** Gemini’s module is `Gromacs` (capital G); lowercase `gromacs` errors on a missing shared dependency. Scripts updated.
- **Config paths under Slurm.** Slurm copies the batch script to a private spool dir, so `$(dirname $0)/configs` resolved into the spool dir and `ions.mdp` was not found. Scripts now anchor paths to `$SLURM_SUBMIT_DIR`; **always submit from the repo root.**
- **GPU bonded offload on EM.** `-bonded gpu` requires a dynamical integrator; EM uses `steep`, so the EM `mdrun` now uses `-nb gpu` only. NVT/NPT/ production use `md` and keep `-nb/-pme/-bonded gpu` (audited).

Also removed the obsolete `ns_type` entry from all `mdp` files (ignored under the Verlet scheme).

Caveat: Gemini scratch is not durable — move logs/outputs off scratch after each run (an earlier prep’s logs were lost to a purge). Sanity check per run: `pdb2gmw.log` must report **4 disulfides** for HEL.

Expansion to multiple antigens (2026-06-21). Switched the box from dodecahedron to **cubic** (interpretability; uniform across systems) and re-ran HEL for consistency. Staged three further apo-antigen systems from the benchmark, hybrid-sourced (strip antigen chain(s) from the complex; numbering kept aligned to the labels). Provenance:

- **HPr** — 2JEL chain P (res 1–85, no Cys); `md/2JEL/apo`.
- **VEGF** — 1BJ1 chains V+W (res 14–107, disulfide-linked homodimer; `pdb2gmw -merge all` for the inter-chain SS); `md/1BJ1/apo`.
- **Bont/A1 Hc** — 2NYY chain A res 872–1295 (C-terminal receptor-binding domain only; catalytic-domain Zn/Ca excluded; res 872 is a domain cut → charged N-terminus); `md/2NYY/apo`.

All four (HEL+these) submitted overnight as prep→production chains (`sbatch --dependency=afterok`).

Deferred: MT-SP1 (2 missing internal residues), Bont/A2 (homology models), hGH (provisional labels). Per-system disulfide expectations recorded in each `prep.sh`.

Structure repair (2026-06-29). Audited REMARK 465 for the incomplete staged structures and repaired two of three by AlphaFold-loop grafting with restrained closure (method + RMSDs in §11; scripts `analysis/repair_structure.py`, `finalize_structure.py`, `close_loops.py`; host env `env/repair-environment.yml`). **hGH** (1HGU, K38–E39 grafted from AFDB P01241) and **tissue factor** (1AHW, loop V83–G90 from AFDB P13726; AF-crystal global C α -RMSD 1.12 Å) are now frozen as `structures/{hGH_1HGU,TissueFactor_1AHW}_fixed.pdb` — continuous, correct disulfides, clash-free — so both are ready to prep. **HCMV gB** (5C6T) is *absent from AlphaFold DB* (viral) and its 32-residue gap is de-novo prediction; deferred pending an AF3 (AlphaFold Server) model. Net: eight antigens ready to `sbatch` on Gemini.

Dynamics-feature push (2026-06-30 / 07-01) — what happened, including the wrong turns. Two days of apo-MD feature extraction and scoring against the *energy* label (alanine-scan $\Delta\Delta G_{\text{bind}}$), on 6 Tier-1 antigens (1AKI 2JEL 1BJ1 1JRH 1AHW 1HGU; ~ 9 antibody-epitopes, ~ 97 scanned residues). Features computed on Gemini: RMSF, trajectory SASA and SASA fluctuation, slow-mode (positional-PCA) participation, DCCM coupling, $\phi/\psi/\chi_1$ dihedral flexibility, per-residue H-bond counts (intra-protein and to water), and gRINN intra-protein interaction energy. Evaluation: *surface-restricted directional AUROC* — restrict to surface residues (trajSASA above the per-antigen median), then $P(\text{feature}_{\text{hotspot}} > \text{feature}_{\text{rest}})$, reporting the sign; leave-one-antigen-out for generalisation. Under it, raw exposure (SASA) collapses to ≈ 0.51 (chance), confirming the baseline the project is built to beat.

Four bugs surfaced, each of which moved a headline number: (i) the alanine-scan $\Delta\Delta G_{\text{bind}}$ table is in antibody-complex PDB numbering while the apo-MD structures are numbered differently — every feature $\leftrightarrow \Delta\Delta G_{\text{bind}}$ join on raw residue number mislabelled 4 of 6 systems (only 39/101 labelled residues matched by wild-type identity; 101/101 after per-system offsets 1AHW–3, 1BJ1–13, 1HGU–1, 1JRH–10; persisted in `benchmark/labels/numbering_offsets.json` + `ddg_by_md_residue.csv`, WT-verified). (ii) gRINN’s output is 0-indexed while every other file is 1-indexed — the energy join was off by one residue; correcting it (`grinn.resid + = 1`) **collapsed interaction-energy detection** 0.67 $\rightarrow$ 0.57: the apparent energy signal, and the “internal-energy / protrusion / under-satisfied side-chain” mechanism built on it, were an alignment artefact — **that narrative is dead**. (iii) gRINN for VEGF ran on a single protomer of the obligate homodimer (1BJ1, 94 vs 188 residues) — energy scale and PC1 broken; 1BJ1 excluded. (iv) reproducibility gaps: the reassessment computation lived only in an ephemeral session (committed script was a stub), the two figures were not regenerable, and the committed trajSASA carried NaN for UNK residues. A runnable evaluation script now replaces the stub (`analysis/reassess_corrected_labels.py`): it applies the offsets with a WT-identity assertion, fixes the gRINN index in one place, and locks the numbers.

Corrected, reconciled headline (5 systems, 1BJ1 dropped; surface, per-antigen AUROC): H-bonds to water 0.66 (best single feature, survives every correction), H-bonds intra 0.60, trajectory SASA 0.57, interaction energy (vdW) 0.57, DCCM 0.56, backbone flexibility 0.52, PC1 participation 0.50, SASA fluctuation 0.40. The HB-water *gradient* among labelled residues is $\rho = -0.33$ (an earlier -0.50 was wrongly surface-restricted, dropping 2JEL, which runs the other way, $\rho = +0.27$); it is not unanimous. **Honest state:** no single antigen-only feature is a strong predictor of the energy label at $N \approx 6$; every exposure/flexibility feature is weak (as the thesis predicts), and interaction energy is now dead. The real bottleneck is N — no transferable linear combination beats the best single feature (LOAO logistic goes below chance) — which motivates scaling the label set. *Still open:* reconcile the HB-water detection number between the runnable script (≈ 0.52 on the shared surface set) and the Science-kernel 0.656 (a residual AUROC/negative-set difference); regenerate

the feature figures once locked; the FRESEAN low-frequency-mode run (proper replacement for the PCA placeholder) is in progress. *Note for a later revision:* the draft energetics narrative rests on the now-retracted energy result and needs rewriting once FRESEAN and the label scale-up land — flagged, not yet edited.

References

- [1] Randi Vita, Swapnil Mahajan, James A. Overton, Sandeep Kumar Dhanda, Sheridan Martini, Jason R. Cantrell, Daniel K. Wheeler, Alessandro Sette, and Bjoern Peters. The immune epitope database (IEDB): 2018 update. *Nucleic Acids Research*, 47(D1):D339–D343, 2019.
- [2] E. Westhof, D. Altschuh, D. Moras, A. C. Bloomer, A. Mondragon, A. Klug, and M. H. V. Van Regenmortel. Correlation between segmental mobility and the location of antigenic determinants in proteins. *Nature*, 311:123–126, 1984.
- [3] Stefano A. Serapian, Filippo Marchetti, Alice Triveri, Giulia Morra, Massimiliano Meli, Elisabetta Moroni, Giuseppe A. Sautto, Andrea Rasola, and Giorgio Colombo. The answer lies in the energy: How simple atomistic molecular dynamics simulations may hold the key to epitope prediction on the fully glycosylated SARS-CoV-2 spike protein. *The Journal of Physical Chemistry Letters*, 11(19):8084–8093, 2020.
- [4] Sarah Sirin, James R. Apgar, Eric M. Bennett, and Amy E. Keating. Ab-bind: Antibody binding mutational database for computational affinity predictions. *Protein Science*, 25(2):393–409, 2016.
- [5] Lei Jin, Brian M. Fendly, and James A. Wells. High resolution functional analysis of antibody-antigen interactions. *Journal of Molecular Biology*, 226(3):851–865, 1992.
- [6] Justina Jankauskaitė, Brian Jiménez-García, Justas Dapkūnas, Juan Fernández-Recio, and Iain H. Moal. SKEMPI 2.0: an updated benchmark of changes in protein–protein binding energy, kinetics and thermodynamics upon mutation. *Bioinformatics*, 35(3):462–469, 2019.
- [7] Yili Li, Hongmin Li, Sandra J. Smith-Gill, and Roy A. Mariuzza. Three-dimensional structures of the free and antigen-bound Fab from monoclonal antilysozyme antibody HyHEL-63. *Biochemistry*, 39(21):6296–6309, 2000.
- [8] T. N. Bhat, G. A. Bentley, G. Boulot, M. I. Greene, D. Tello, W. Dall’Acqua, H. Souchon, F. P. Schwarz, R. A. Mariuzza, and R. J. Poljak. Bound water molecules and conformational stabilization help mediate an antigen-antibody association. *Proceedings of the National Academy of Sciences*, 91(3):1089–1093, 1994.
- [9] E. A. Padlan, E. W. Silverton, S. Sheriff, G. H. Cohen, S. J. Smith-Gill, and D. R. Davies. Structure of an antibody-antigen complex: crystal structure of the HyHEL-10 Fab-lysozyme complex. *Proceedings of the National Academy of Sciences*, 86(15):5938–5942, 1989.
- [10] Lata Prasad, E. Bruce Waygood, Jeremy S. Lee, and Louis T. J. Delbaere. The 2.5 Å resolution structure of the Jel42 Fab fragment/HPr complex. *Journal of Molecular Biology*, 280(5):829–845, 1998.
- [11] Yves A. Muller, Yvonne Chen, Hans W. Christinger, Bing Li, Brian C. Cunningham, Henry B. Lowman, and Abraham M. de Vos. VEGF and the Fab fragment of a humanized neutralizing

- p>antibody: crystal structure of the complex at 2.4 Å resolution and mutational analysis of the interface.
- Structure*
- , 6(9):1153–1167, 1998.
- [12] Yvonne Chen, Christian Wiesmann, Germaine Fuh, Bing Li, Hans W. Christinger, Patrick McKay, Abraham M. de Vos, and Henry B. Lowman. Selection and analysis of an optimized anti-VEGF antibody: crystal structure of an affinity-matured Fab in complex with antigen. *Journal of Molecular Biology*, 293(4):865–881, 1999.
 - [13] Christopher J. Farady, Pascal F. Egea, Eric L. Schneider, Molly R. Darragh, and Charles S. Craik. Structure of an Fab–protease complex reveals a highly specific non-canonical mechanism of inhibition. *Journal of Molecular Biology*, 380(2):351–360, 2008.
 - [14] Eric L. Schneider, Melody S. Lee, Aida Baharuddin, David H. Goetz, Christopher J. Farady, Mick Ward, Cheng-I Wang, and Charles S. Craik. A reverse binding motif that contributes to specific protease inhibition by antibodies. *Journal of Molecular Biology*, 415(4):699–715, 2012.
 - [15] Consuelo Garcia-Rodriguez, Raphael Levy, Joseph W. Arndt, Charles M. Forsyth, Ali Razai, Jianlong Lou, Isin Geren, Raymond C. Stevens, and James D. Marks. Molecular evolution of antibody cross-reactivity for two subtypes of type A botulinum neurotoxin. *Nature Biotechnology*, 25(1):107–116, 2007.
 - [16] Klaus Hofstädter, Fiona Stuart, Luyong Jiang, Jan W. Vrijbloed, and John A. Robinson. On the importance of being aromatic at an antibody-protein antigen interface: mutagenesis of the extracellular interferon γ receptor and recognition by the neutralizing antibody A6. *Journal of Molecular Biology*, 285(2):805–815, 1999.
 - [17] Mingdong Huang, Rashid Syed, Enrico A. Stura, Martin J. Stone, Randy S. Stefanko, Wolfram Ruf, Thomas S. Edgington, and Ian A. Wilson. The mechanism of an inhibitory antibody on TF-initiated blood coagulation revealed by the crystal structures of human tissue factor, Fab 5G9 and TF·5G9 complex. *Journal of Molecular Biology*, 275(5):873–894, 1998.
 - [18] Sumana Chandramouli, Claudio Ciferri, Pavel A. Nikitin, Stefano Caló, Rachel Gerrein, Kara Balabanis, James Monroe, Christy Hebner, Anders E. Lilja, Ethan C. Settembre, and Andrea Carfi. Structure of HCMV glycoprotein B in the postfusion conformation bound to a neutralizing human antibody. *Nature Communications*, 6:8176, 2015.
 - [19] Edgar Davidson and Benjamin J. Doranz. A high-throughput shotgun mutagenesis approach to mapping B-cell antibody epitopes. *Immunology*, 143(1):13–20, 2014.
 - [20] Matthew Z. Tien, Austin G. Meyer, Dariya K. Sydykova, Stephanie J. Spielman, and Claus O. Wilke. Maximum allowed solvent accessibilities of residues in proteins. *PLoS ONE*, 8(11):e80635, 2013.